


null-drop depth additivity — GO

Each affected token drops one expert (lowest mass) per layer, no renormalization; same affected sample across L. Additive-residual hypothesis predicts $KL \sim L$. Monotone: True; super-linear: False.

Is null-drop cost additive in depth?



Routing/top-1 stability vs depth

